

Novatech Revenue Intelligence Dashboard

Report for Sarah Chen, VP of Revenue

Prepared by Rhyah Austrie S. Reyes, Business Intelligence Analyst

Background

This dashboard pulls together NovaTech's three systems, which consist of CRM deals, marketing campaigns, and support tickets, into one place. It covers closed deals, marketing leads, and support tickets on record, connected through a shared account ID. The goal was to stop the Monday morning routine of pulling numbers from three different tools and give the team one dashboard they can check on their own, including a chat feature that answers questions.

Data Strategy

The three systems don't store data the same way, so the first step was to clean each one individually. Dates that were stored as plain text got converted to real date fields. Money fields got converted to decimals instead of text. I also added three new calculated fields: days to close (time from opening to closing), campaign ROI (revenue vs. spend), and ticket resolution time in hours.

To connect the three sources, I used the CRM deals table as the base, since a closed deal is the core unit of revenue. I joined marketing and support to it using a left join, which keeps every CRM deal even if that account never generated a lead or filed a ticket. I chose this over an inner join, which would have dropped deals with no matching marketing or support record and made revenue look smaller than it is.

One thing worth flagging is that joining to two tables at once causes row multiplication. Since one account can have many leads and many tickets, the combined table ends up with 63,420 rows even though there are only 499 deals underneath. This isn't a mistake, but just how joins work when one account has multiple matches. It does mean the combined table should only be used to look at things at the account level, like on the Customer Health sheet. For anything involving deal revenue, the Sales Pipeline sheet reads straight from the clean CRM table instead, so a deal's value never gets counted more than once.

Dashboard Design Rationale

The dashboard has three sheets that follow the path a customer actually takes: Marketing Funnel, Sales Pipeline, and Customer Health. Each sheet opens with KPI cards so you can get the number at a glance, then moves into charts and tables for anyone who wants to dig further. Filters for campaign, channel, segment, region, and product also let the team answer follow-up questions instead of coming back to me for a new report every time.

On the Sales Pipeline sheet, clicking a point on the deal closing time chart filters every other chart on that sheet to the same period, so you can click into a specific month and see everything update at once.

On the Customer Health sheet, clicking an account in the at-risk accounts table takes you straight to that account's deal history on the Sales Pipeline sheet, which was the main thing you asked to connect support and sales for the same account.

How Topic Configuration Affected Q&A Accuracy

A Topic is what teaches Quick's chat assistant, which I'll call it Q, how to read NovaTech's specific terms and data. Before I set one up, a revenue question came back with \$100,900,682. The real number is \$707,201. Q pulled from the combined dataset described above, the one affected by row multiplication, instead of the clean CRM numbers.

After I configured the Topic, explaining that revenue means deal value on Won deals only, that deals and tickets should be counted rather than summed, and that a real conversion rate needs total leads per channel, not just deal value, the same kind of question came back correct.

When I asked the exact same question again, *"What is total revenue from closed-won deals?"*, Q answered \$707,201, which matches the KPI on the dashboard. It also pointed out on its own that this number was different from the one it gave earlier, without me asking it to.

Two things still need work, and I've written them down in the Q Exploration Log with a suggested fix. First, if you just ask for average deal size without saying Won deals, Q averages in the Lost deals, which are all worth \$0, and that drags the number down. Second, we couldn't answer a question about Enterprise accounts because the customer tier field isn't included in the Topic yet, even though it exists in the data. Both are quick fixes to the Topic, not data problems.

Key Insights and Recommended Actions

- Partner Referral brought in 60.9 percent of closed-won deals, 251 of 412, more than every other channel combined. Organic Search brought in the most leads, 2,673, but converted the worst, at 47.89 percent.
 - Marketing should put more budget behind referral incentives and take a hard look at what's happening with Organic Search before the next budget cycle.
- Total revenue from closed-won deals is \$707,201 across 315 deals, which is 63 percent of all 499 deals in the system. Average deal size for Won deals is \$2,245.08.
 - Sales leadership should pull the loss reasons on the 184 Lost deals and see if there's a pattern worth fixing.
- Average days to close dropped sharply over the period we have data for, from a high of 191 days down to as little as 1 day on the most recent deals. Whatever changed, it's working.
 - Sales Ops should figure out what happened on the faster deals and see if it can be repeated across the rest of the team.
- Five accounts, ACCT 007, ACCT 039, ACCT 053, ACCT 070, and ACCT 083, show \$0 in deal value but still have open support tickets, ranging from 18 to 960 tickets each.

They're either churned accounts still eating up support time, or something's off in how they're recorded. Separately, ACCT 041, YieldMax Software, has 5,010 support tickets against a deal worth \$13,601,148, a very high support load relative to its size.

- Customer Success should review all six accounts within the next two weeks, the five zero-value accounts for reclassification, and ACCT 041 for an account health check.
- Resolution time barely changes by priority level. Low-priority tickets take about 61 hours, and high-priority tickets take about 60. If priority isn't actually speeding anything up, the triage process probably needs a second look.

Where AI Analysis Agreed or Disagreed with the Dashboard

Once I phrased questions to point at the right dataset and the right deal stage, Q's answers lined up with the dashboard exactly, on total revenue and on the Won versus Lost split. That's a good sign the numbers on the dashboard hold up under a different kind of query. Q also helped surface the at risk account list directly from the dashboard's own table, which is how the zero-value accounts and the high ticket volume on ACCT 041 were confirmed against the source data rather than estimated.

Where Q and the dashboard didn't match, it always came back to how the question was phrased or a gap in the Topic setup, not an actual disagreement about the data. An unscoped revenue question pulled from the wrong table. An unscoped deal size question included lost deals it shouldn't have. A question about Enterprise accounts hit a field that wasn't in the Topic yet. A request for the full ten account list came back truncated to two rows, a display limit rather than a data problem. All of these are documented in the [Q Exploration Log](#) along with what would fix them.

Recommended Next Steps

- Add customer tier to the Topic, with a note on what Basic, Enterprise, and Plus mean, so Q can filter by account tier correctly.
- Set a rule, or a default filter, that revenue and deal size questions mean Won deals only, so the numbers stay consistent without someone having to specify it every time.
- Have Customer Success review the five zero-value accounts and ACCT 041 within the next two weeks.
- Have Support revisit how tickets get triaged, since priority level doesn't seem to be changing resolution speed right now.